


Build an operating system for your AI agents.

Give Codex, Claude Code, and agentic tools intake, truth, proof, approval, and learning.



AI coding tools are becoming teammates.

The risk is not that agents help. The risk is that nobody manages the system they work inside.

THE MISTAKE

One big prompt is not an operating model.

Loose prompt

No owner, no repo lane, no proof, no approval gate.

Tool hopping

Codex, Claude Code, browser work, and app spaces all mixed together.

Silent changes

Work happens without a human-visible trail.

Autonomy theater

The system looks busy but cannot prove business movement.

THE MODEL

Give each layer a job.

Operating spaces

Channels, projects, app spaces, and threads for routing.

Source truth

Repos, docs, design files, checks, and decision records.

Codex

Synthesis, repo work, reports, and approval packets.

Specialist tools

Claude Code, ChatGPT, browser sessions, and focused build tools.

Runtime later

Hermes identity and dispatch only after gates are proven.

Human

Approves public, financial, production, and customer-facing action.



Use a repeatable execution loop.

1

Brief the agent with outcome, source truth, constraints, and proof required.

2

Work in a branch, thread, project, or bounded app workspace.

3

Run checks and inspect the actual artifact.

4

Return proof with links, decision needed, and learning captured.

Agents can prepare. Humans approve irreversible moves.

Allowed by default

Research, draft, inspect, test, summarize, prepare branches, produce reports, create approval packets.

Approval required

Publish, merge, deploy, spend, change access, message customers, transfer domains, schedule social posts.

THE PROMPT

Give agents the operating packet, not a vibe.

A Objective: what should exist when done?

B Context: repo, app space, docs, brand, sources.

C Boundaries: what must not happen?

D Proof: what link, test, file, or screenshot closes the loop?

WHY THIS MATTERS NOW

The labs are moving toward longer-running, governed AI work.

OpenAI Codex

Official Codex pages frame coding agents across app, editor, terminal, cloud, skills, and automations.

Long-running work

OpenAI's 2026 Codex guide emphasizes continuity, verifiable steps, and human oversight.

Claude Code

Anthropic documents Claude Code as an agentic coding tool across terminal, IDE, desktop, and browser.

Enterprise controls

OpenAI's 2026 enterprise update emphasizes usage analytics, spend controls, and admin visibility.

Start with five operating spaces.

1

Status: what changed, what matters, what needs approval.

2

Intake: requests before assignment.

3

Execution: active work only.

4

Truth: repo, docs, design, domain, and risk evidence.

5

Approval: public-facing and irreversible gates.

SAVE THIS

Agents need an operating system.

Not more random prompts. Not more disconnected tools. A real system of lanes, proof, approvals, and review.

Sources: OpenAI Codex; OpenAI Codex long-running work guide; OpenAI enterprise controls; Anthropic Claude Code overview.