

Your AI agents do not need more prompts.

They need a workbench: a place for work to enter, proof to form, and risky moves to wait for human approval.

WORK
BENCH



Intake

what changed
and why



Bench

bounded
active work



Source

repo, docs,
proof



Proof

checks and
evidence



Approve

human gate



Learn

carry forward

Prompts alone create drift.

When every request starts from a blank chat, the agent has no lane, no source of truth, no proof standard, and no stop condition.

Loose ask

No owner, scope, boundary, or proof requirement.

Tool hopping

Work spreads across chats, branches, docs, and apps.

Silent change


Something moves, but nobody knows what changed.

Autonomy theater

It looks advanced until approval, evidence, or source truth is needed.

A workbench, not a chat.

Give each agent workflow a physical operating surface: intake, bench, source, proof, approval, and learning.



1

Requests enter before agents act.

Capture objective, owner, and risk.

2

Work happens inside a bounded space.

Branch, thread, project, repo, or app workspace.

3

Proof returns before approval.

Links, checks, screenshots, source notes, and decision needed.

4

Learning changes the next run.

Keep what worked. Cut what failed.

Every agent workflow needs a place to live.

I

Intake

Requests, context, owner, priority, and risk.

B

Bench

Branch, thread, project, or bounded app workspace.

S

Source

Repo, docs, product truth, citations, and constraints.

P

Proof

Tests, screenshots, diffs, citations, and QA notes.

A

Approval

Human gate for public, production, money, and customer risk.

L

Learning

Decision notes, performance, corrections, and next rule.

Agents prepare. Humans approve.

Allowed by default

- Research and summarize
- Draft copy or code
- Run local checks
- Prepare screenshots
- Return proof packets

Needs a human yes

- Publish or schedule
- Deploy production
- Merge to main
- Message customers
- Spend or change access



Give the agent the work order, not a vibe.

Operating packet

USE FIRST

Objective	What should exist when this is done?
Context	Repo, docs, brand, customer, sources, constraints.
Boundary	Allowed files, non-goals, stop conditions.
Checks	Build, lint, test, visual QA, source review.
Proof	Links, screenshots, diffs, source list, decision needed.
Approval	What needs a human yes before action?

A confident agent is not a trustworthy system.

Trust arrives when the agent returns evidence another human can inspect.

No proof, no public move.

The proof packet is the difference between useful delegation and blind automation.

1 Source links or repo paths

2 Checks run and result

3 Screenshot or rendered asset

4 Claim-risk label

5 Human decision requested

Capability is rising. Governance has to rise with it.

Official product docs frame coding agents as systems that can work with codebases, tools, and longer software tasks. Treat that as a reason to improve the operating model.

OPENAI

Codex

Official Codex materials describe a cloud-based coding agent for software work.

openai.com/codex/

OPENAI DOCS

Codex cloud

Cloud docs frame Codex as a way to delegate software engineering tasks in a controlled environment.

developers.openai.com/codex/cloud

ANTHROPIC

Claude Code

Claude Code docs describe an agentic coding tool for terminal, IDE, browser, and cloud workflows.

code.claude.com/docs/en/overview

Build the workbench before you add autonomy.

1	Name the intake lane.	Where requests arrive before assignment.
2	Pick the active workspace.	Branch, thread, project, repo, app, or room.
3	Define source truth.	Repos, docs, specs, brand packs, and citations.
4	Require proof.	Rendered assets, tests, checks, links, and screenshots.
5	Write the approval gate.	Public, production, financial, customer, access.
6	Save the learning.	What to make, cut, and carry forward.

The real stack is not more tools.

It is intake, workspace, truth, proof, approval, and learning.

**Agents prepare the work.
Humans approve irreversible moves.**

Comment "agent OS" for the setup template.

Sources checked: OpenAI Codex, OpenAI Codex cloud docs, Anthropic Claude Code overview. No public publish or schedule action is approved by this deck.

Intake

Bench

Source

Proof

Approval

Learning